

# MPI: Project - Game of Life

You are going to write an MPI parallel version of the Game of Life!

*"Game of Life" simulation by John Conway*



2D array of cells.

Each cell can have one of two possible states: **alive** or **dead**

Initialise with random states and then evolve according to rules:

- If exactly 3 neighbours are alive, cell will be alive (if already alive, remains alive; if was dead, becomes alive)
- If exactly 2 neighbours are alive, no change in cell status.
- All other cases, cell is dead (if was dead, remains dead; if was alive, becomes dead).

Assume **periodic boundary conditions** (don't forget you need diagonals)



# MPI: Project - Game of Life

I suggest you proceed with the following steps:

1. Write an outline of the steps needed to compute the game
2. Write a serial version of the game. To debug, use a small grid (e.g. 4x4) and a single step. You will need this code to check the correctness of your parallel version and to compare speeds. Keep a track (print out) of the total number of alive cells -- this can be your simple metric for whether things are working or not. I suggest you use ghost cells for the boundary conditions.
3. Add the MPI initialisation and finalisation routines to your code and call it a parallel version. Doing this makes a code that runs the same program redundantly on all the processors you specify. Print the rank and the metric to make sure all is working.
4. Do a domain decomposition of the problem to parallelise the problem truly. Since you are programming in Fortran, I recommend initially dividing by columns. This is because of the way arrays are contiguous in memory (i.e. column major in Fortran). Use ghost cells for interior boundaries. Debug with a partition of 2.
5. Mega version: Divide the domain the OTHER way. This means you will need to use derived data types, since these are not the natural ways for the language.
6. Giga version: Decompose the problem in both directions.
7. Tera version: Use a virtual topology. Use `MPI_CART_CREATE`, and use `MPI_CART_SHIFT` to create the ghost cells.
8. Peta version: Use parallel I/O to write out the solution grid and visualise it. There will be extra credit for any cute ways of visualizing ☺

Pass level


100% level

# MPI: Project - Game of Life

Report: due Wed June 8<sup>th</sup> at 11:59pm

Write a SHORT report that

1. Describes the PCAM process that you went through to design your code (can be BRIEF: <0.5 page)
2. Documents everything your code needs to run so that I can run it and check it on grape. Include a hardcopy of your code.
3. To prove that your code works in parallel, start with a 20x20 grid and initialise the grid with the configuration shown here (to the right) at the top left of your grid( and nothing elsewhere). Run the code for 4 steps and notice that this returns back to the same pattern but with everything shifted diagonally by 1 square. Now run the code for 80 steps. The pattern should return to where it started. Show the grid at steps 0,20,40,80 when the code is run on 4 processors.
4. Come up with a simple performance model and test it out. Remember, you have the communication time constants ( $t_s$  and  $t_w$ ) from your latency program (earlier HW) and the computation time ( $t_c$ ) can be calculated from your ones.f90 (earlier HW) too.



# MPI: Project - Game of Life

## To submit your project:

1. Make a directory wherever you are working called `AM250_project_<your_lastname>`
2. Copy ALL the necessary files to run your code(s), including a compiled binary executable, the batch jobfile if necessary, etc, to this directory and include a `README.TXT` that describes (very briefly) the commands necessary to compile and run your code, e.g.

“On grape, do the following to compile and run my code:

To change the grid size and the number of processors, edit `my_fortran.f90` and change `ngrid` and `nprocs`

```
mpif90 -o my_exec my_fortran.f90
```

```
qsub my_batchfile.cmd or mpirun -np 4 my_exec
```

The output data will be in `my_datafile`

To look at the data, type `more my_datafile`”

3. Copy your report in PDF form into your user name directory too. Call it `report_username.pdf`
4. Tar the directory into a single tarfile named by your last name:

```
cd .. (if you start from your code directory)
```

```
tar cvf <your_lastname>.tar AM250_project_<your_lastname>
```

5. Submit the tarfile to the Canvas page, AND SUBMIT THE REPORT AS A SEPARATE FILE TOO.